

LSSO: Learnable Low-Rank Sylvester Solves for Efficient Bidirectional Encoders

Anonymous Authors

Abstract

We introduce LSSO, a token mixer for bidirectional encoders built around a learned low-rank linear solve. Rather than routing values through an explicit pairwise attention matrix, LSSO constructs an input-dependent relation field and solves

$$(\mu I + \gamma U(X)U(X)^\top)Y = C(X).$$

This is the left-sided case of a Sylvester equation. The positive shift makes the system well posed, while the Woodbury identity reduces the computation to an $r \times r$ solve. For fixed relation features, the operator has a bounded spectral response and minimizes a strongly convex correction energy. Across retrieval, MS MARCO-to-BEIR transfer, and vision classification, LSSO matches or approaches attention-based encoders while using substantially fewer mixer MACs. Ablations, rank pruning, and sequence-scaling experiments further show when the global correction is useful and where the current formulation falls short.

1 Introduction

Self-attention [21] is the standard token mixer in modern bidirectional encoders. Given query, key, and value projections, it computes

$$Y_{\text{attn}} = \text{softmax}(QK^\top / \sqrt{d})V,$$

and lets every token route information from every other token. That flexibility comes with quadratic score and value-routing costs. Most efficient attention methods reduce the cost by approximating or sparsifying the routing matrix. LSSO asks a different question: can a global token interaction be learned without constructing such a matrix at all?

For an input sequence $X \in \mathbb{R}^{B \times N \times D}$, LSSO learns a low-rank relation feature $U(X)$ and a content state $C(X)$. For each head it solves

$$(\mu I_N + \gamma U U^\top)Y = C, \quad \mu > 0, \quad \gamma \geq 0,$$

and projects the result back to the model dimension. Although $U U^\top$ is positive semidefinite, its entries $u_i^\top u_j$ may be positive or negative. The system can therefore encode signed token relations while remaining invertible by construction.

The guiding interpretation is:

LSSO = isotropic content pass-through
– low-rank global correction.

Attention learns a row-normalized routing rule; LSSO learns the relation features of a positive-definite inverse filter. The two mixers therefore have different inductive biases even when their input and output shapes match.

The main contributions are:

1. We formulate bidirectional token mixing as a learned, positive-shifted low-rank solve whose implementation only solves an $r \times r$ system.
2. We characterize the fixed- U operator through its spectrum, conditioning, norm bound, and equivalent convex energy.
3. We evaluate the mixer on retrieval, zero-shot transfer, vision, rank pruning, diagnostics, and long-sequence scaling, including settings where LSSO does not outperform attention.

2 Related Work

Efficient attention. Linformer [23], Performer [5], Nyströmformer [25], Longformer [3], and BigBird [27] reduce the cost of self-attention through low-rank projections, kernel approximations, landmark approximations, or sparse attention patterns. Despite these differences, they retain attention as the underlying interaction rule. LSSO instead replaces that rule with an input-dependent inverse filter.

Token mixers beyond attention. MLP-Mixer [19], gMLP [14], FNet [13], Hyena [17], and state-space models such as Mamba [11] demonstrate that useful token mixing does not require dense attention. LSSO explores a complementary construction: an implicit low-rank solve designed specifically for non-causal encoders.

Implicit layers and learned solves. Deep equilibrium models [1], implicit layers, graph smoothing, and regularized least-squares methods place linear or nonlinear solves inside neural networks. LSSO is narrower than a general equilibrium layer: it applies one input-conditioned linear solve with a deliberately constrained positive-definite matrix.

3 Method

3.1 Encoder Block

We use LSSO in the token-mixer position of a pre-normalized encoder block:

$$X' = X + \text{LSSO}(\text{LN}(X)), \quad X_{\text{out}} = X' + \text{MLP}(\text{LN}(X')).$$

In controlled comparisons we change only this sublayer; normalization, residual paths, feed-forward layers, positional embeddings, and task heads are shared across models.

3.2 LSSO Definition

For each head, let

$$U = XW_u \in \mathbb{R}^{B \times N \times r}, \quad C = XW_c \in \mathbb{R}^{B \times N \times d_h}.$$

We optionally RMS-normalize U along its rank dimension and solve

$$(\mu I_N + \gamma U U^\top) Y = C, \quad (1)$$

where

$$\mu = \text{softplus}(\theta_\mu) + \epsilon, \quad \gamma = \gamma_{\max} \sigma(\theta_\gamma).$$

The default experiments use $\gamma_{\max} = 0.3$ and $\theta_\gamma = -4$, which gives an initial γ/μ of roughly 7.8×10^{-3} when $\theta_\mu = 0$. More conservative settings such as $\gamma_{\max} = 0.1, \theta_\gamma = -6$ make the global correction nearly vanish at initialization; in preliminary runs they consistently underused the global term.

The head outputs are concatenated and projected:

$$\text{LSSO}(X) = YW_o.$$

3.3 Woodbury Implementation

The $N \times N$ matrix in Eq. 1 never needs to be formed or inverted. Applying the Woodbury identity [24] gives

$$Y = \mu^{-1} C - \mu^{-1} U (I_r + \gamma \mu^{-1} U^\top U)^{-1} (\gamma \mu^{-1} U^\top C). \quad (2)$$

This expression remains valid at $\gamma = 0$ and requires only an $r \times r$ solve. The core solve terms cost

$$O(Nr^2 + Nr d_h + r^3)$$

per head. Including projections, the mixer-only MAC formula used in our tables is

$$ND(Hr + D) + ND^2 + HNr^2 + 2NrD + Hr^3.$$

For MHA we count QKV projections, attention scores, attention-value product, and output projection:

$$4ND^2 + 2N^2D.$$

4 Theoretical Properties

We first study the linear map induced by a fixed U . These results describe the operator itself; they do not establish convergence of end-to-end training.

Proposition 1 (Well-posedness). *Let $A = \mu I_N + \gamma U U^\top$ with $\mu > 0$ and $\gamma \geq 0$. Then A is symmetric positive definite. Therefore, for any C , the system $AY = C$ has a unique solution.*

Proof. For any nonzero z ,

$$z^\top A z = \mu \|z\|_2^2 + \gamma \|U^\top z\|_2^2 \geq \mu \|z\|_2^2 > 0.$$

□

Proposition 2 (Spectral shrinkage). *Let $U U^\top = Q \Lambda Q^\top$ be an eigendecomposition over the nonzero eigen-directions of $U U^\top$, with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_k)$ and $\lambda_i > 0$. Let $P_\perp = I - Q Q^\top$. Then*

$$A^{-1} = Q \text{diag}\left(\frac{1}{\mu + \gamma \lambda_i}\right) Q^\top + \mu^{-1} P_\perp.$$

Thus LSSO applies an inverse spectral filter on the span of U and an isotropic scale on its orthogonal complement. This gives a more precise view of the distinction from row-normalized attention.

Proposition 3 (Bounded fixed- U operator norm). *For fixed U ,*

$$\|\text{LSSO}(C, U)\|_F \leq \mu^{-1} \|C\|_F.$$

With an output projection,

$$\|W_o \text{LSSO}(C, U)\|_F \leq \|W_o\|_2 \mu^{-1} \|C\|_F.$$

Proof. Since $A \succeq \mu I$, $\lambda_{\min}(A) \geq \mu$ and $\|A^{-1}\|_2 \leq \mu^{-1}$. The Frobenius bound follows from $Y = A^{-1}C$. □

Proposition 4 (Conditioning control). *For $A = \mu I + \gamma U U^\top$,*

$$\kappa_2(A) \leq 1 + \frac{\gamma}{\mu} \|U\|_2^2.$$

These bounds motivate the diagnostics used in our experiments. The ratio γ/μ directly scales the possible condition-number increase; effective rank measures how much of the rank subspace is used; and correction ratio measures how strongly the low-rank correction participates in the output. RMS normalization of U helps control the scale of U , although it is not by itself a dimension-free spectral-norm guarantee.

Theorem 1 (Energy minimization equivalence). *For $\mu > 0$ and $\gamma \geq 0$, the LSSO output is the unique minimizer of*

$$\mathcal{E}(Y) = \frac{\mu}{2} \|Y\|_F^2 + \frac{\gamma}{2} \|U^\top Y\|_F^2 - \langle C, Y \rangle.$$

Proof. The Hessian of $\mathcal{E}$ is $\mu I + \gamma U U^\top$, which is positive definite by the well-posedness proposition. Setting the gradient to zero gives $(\mu I + \gamma U U^\top)Y = C$. □

Equivalently, up to an additive constant,

$$\mathcal{E}(Y) = \frac{\mu}{2} \|Y - \mu^{-1}C\|_F^2 + \frac{\gamma}{2} \|U^\top Y\|_F^2.$$

The content projection supplies the local state $\mu^{-1}C$, and the second term adjusts that state using sequence-wide information.

Proposition 5 (No-global degeneration). *If $\gamma = 0$, then*

$$\text{LSSO}(C, U) = \mu^{-1}C.$$

No token mixing occurs except for pointwise projections.

5 Experiments

5.1 Setup

All main retrieval and vision comparisons use randomly initialized encoders. The retrieval backbone follows a BERT-style bidirectional encoder layout [9]. We reuse standard tokenization but train every encoder weight from scratch. Models are optimized with AdamW [15]. LSSO uses $\gamma_{\max} = 0.3$ and $\theta_\gamma = -4$ unless stated otherwise. Because the feed-forward network is identical across models, tables report mixer-only MACs and isolate the cost of the component being replaced.

5.2 Retrieval Main Results

Table 1 reports 3-seed retrieval results on FIQA [16], NFCorpus [4], and SciFact [22], using their BEIR retrieval formulations [18]. The baseline is a BERT-style MHA encoder, and Nyströmformer uses the authors’ attention module. LSSO-r16 and LSSO-r32 use the same model dimension, depth, number of heads, optimizer, and training schedule.

Table 1: Retrieval main results. Models are randomly initialized BERT-style encoders with $D = 256$, depth 8, heads 8, and document length 512. MACs are mixer-only document-side MACs.

Dataset	Model	Params (M)	Mixer MACs (G)	R@10	MRR@10
FIQA	MHA	14.33	2.147	0.2145 $\pm$ 0.0135	0.0987 $\pm$ 0.0032
FIQA	Nyströmformer	14.33	1.301	0.2428 $\pm$ 0.0039	0.1190 $\pm$ 0.0071
FIQA	LSSO-r16	13.54	0.713	0.2258 $\pm$ 0.0125	0.1126 $\pm$ 0.0061
FIQA	LSSO-r32	13.80	0.908	0.2387 $\pm$ 0.0183	0.1150 $\pm$ 0.0035
NFCorpus	MHA	14.33	2.147	0.5439 $\pm$ 0.0187	0.3809 $\pm$ 0.0145
NFCorpus	Nyströmformer	14.33	1.301	0.5841 $\pm$ 0.0125	0.4323 $\pm$ 0.0132
NFCorpus	LSSO-r16	13.54	0.713	0.5686 $\pm$ 0.0125	0.4121 $\pm$ 0.0036
NFCorpus	LSSO-r32	13.80	0.908	0.5841 $\pm$ 0.0036	0.4333 $\pm$ 0.0220
SciFact	MHA	14.33	2.147	0.6789 $\pm$ 0.0190	0.5994 $\pm$ 0.0138
SciFact	Nyströmformer	14.33	1.301	0.7267 $\pm$ 0.0133	0.6313 $\pm$ 0.0033
SciFact	LSSO-r16	13.54	0.713	0.7022 $\pm$ 0.0107	0.6214 $\pm$ 0.0119
SciFact	LSSO-r32	13.80	0.908	0.7089 $\pm$ 0.0117	0.6247 $\pm$ 0.0080

Relative to MHA, LSSO-r16 removes 66.8% of mixer MACs and LSSO-r32 removes 57.7%. Both improve the mean MHA result on every dataset in the table. Nyströmformer remains the strongest baseline on several metrics, although in our implementation its lower theoretical MAC count did not translate into higher measured throughput.

5.3 MS MARCO to BEIR Transfer

We next ask whether the retrieval result survives a change of domain. Each model is pretrained once on MS MARCO [2] and evaluated without target-task finetuning on five BEIR datasets [18] covering finance, biomedicine, scientific fact checking, argument retrieval, and COVID literature. Unlike the preceding experiments, this setting does not allow the encoder to adapt to each evaluation collection.

All four models use the same randomly initialized BERT-style architecture, MS MARCO training data, tokenizer, pooling, optimizer, three seeds, and document length 512; only the token mixer changes. For each seed, we first average a metric across the five target datasets and then report the mean and standard deviation of those macro averages. Since total parameter counts are similar (13.53–14.33M), Table 2 emphasizes mixer MACs. The released artifacts contain the corresponding per-dataset and per-seed values.

Both LSSO variants preserve MHA-level transfer at a much lower mixer cost. LSSO-r16 slightly exceeds the MHA macro nDCG@10 (0.22973 vs. 0.22945) while using 66.8% fewer mixer MACs. LSSO-r32 has the strongest macro Recall@100 (0.36367) while using 57.7% fewer mixer MACs. Nyströmformer is also cheaper than MHA, but trails both MHA and LSSO on all three macro metrics.

No single rank wins every metric. LSSO-r16 offers the best nDCG@10–cost trade-off, LSSO-r32 reaches the highest Recall@100, and MHA retains the best MRR@10. Taken together, the two LSSO settings lie on the same transfer frontier as MHA while requiring considerably less mixer computation.

5.4 Ablations and Diagnostics

Table 3 aligns the empirical ablations with the theory. The no-global variant sets $\gamma = 0$, reducing the operator to $\mu^{-1}C$ and removing token mixing. Removing RMS normalization weakens the global correction, especially on SciFact. Smaller ranks expose the capacity of the learned relation subspace.

Table 3: Retrieval ablations. Values are mean $\pm$ standard deviation over three seeds.

Dataset	Variant	R@10	MRR@10	γ/μ	Correction ratio
FIQA	no-global r32	0.2052 $\pm$ 0.0137	0.0911 $\pm$ 0.0019	0.0000	0.0000
FIQA	fixed μ/γ r32	0.2428 $\pm$ 0.0126	0.1136 $\pm$ 0.0028	0.0078	0.2130
FIQA	no U RMS norm r32	0.2160 $\pm$ 0.0101	0.1007 $\pm$ 0.0008	0.0085	0.2050
FIQA	full r8	0.2088 $\pm$ 0.0208	0.0923 $\pm$ 0.0139	0.0100	0.1059
FIQA	full r4	0.2042 $\pm$ 0.0045	0.0979 $\pm$ 0.0036	0.0109	0.0742
SciFact	no-global r32	0.6767 $\pm$ 0.0120	0.5811 $\pm$ 0.0058	0.0000	0.0000
SciFact	fixed μ/γ r32	0.7089 $\pm$ 0.0102	0.6088 $\pm$ 0.0066	0.0078	0.2631
SciFact	no U RMS norm r32	0.6933 $\pm$ 0.0133	0.5972 $\pm$ 0.0150	0.0078	0.0471
SciFact	full r8	0.7067 $\pm$ 0.0173	0.6118 $\pm$ 0.0064	0.0078	0.1149
SciFact	full r4	0.7022 $\pm$ 0.0069	0.6146 $\pm$ 0.0107	0.0078	0.0766

Figure 1 shows the same quantities layer by layer. The correction ratio remains nonzero on every task, so the trained models do not collapse to the local projection. Effective rank varies more substantially: SciFact uses a

Table 2: Zero-shot BEIR transfer after MS MARCO pretraining. Retrieval metrics are five-dataset macro averages, reported as mean $\pm$ standard deviation over three seeds. MACs are mixer-only at document length 512. Best values are bold.

Model	Mixer MACs (G)	nDCG@10	Recall@100	MRR@10
MHA	2.147	0.22945 $\pm$ 0.00208	0.36061 $\pm$ 0.00365	0.30225$\pm$0.00442
Nyströmformer	1.270 (−40.9%)	0.20523 $\pm$ 0.00495	0.33456 $\pm$ 0.00837	0.27115 $\pm$ 0.01046
LSSO-r16	0.713 (−66.8%)	0.22973$\pm$0.00232	0.35889 $\pm$ 0.00226	0.29382 $\pm$ 0.01446
LSSO-r32	0.908 (−57.7%)	0.22940 $\pm$ 0.00337	0.36367$\pm$0.00830	0.29858 $\pm$ 0.00624

larger fraction of the rank-32 subspace than FIQA, NF-Corpus, or CIFAR-100.




Figure 1: Layer-wise LSSO diagnostics from trained LSSO-r32 checkpoints.

5.5 Sequence Scaling

We benchmark a single mixer layer at sequence lengths from 128 to 2048, using batch size 8, $D = 256$, 8 heads, bf16 autocast, and no optimizer update. Figure 2 reports forward latency, forward-plus-backward latency, incremental peak memory, and theoretical mixer MACs.




Figure 2: Sequence scaling benchmark on an NVIDIA RTX 5070 Ti. MACs are theoretical per-sample single-layer mixer MACs; latency and memory use batch 8.

At $N = 2048$, LSSO-r16 uses 0.357G mixer MACs compared with 2.684G for MHA, an 86.7% reduction, and reduces forward+backward latency from 2.61 ms to 1.88 ms. LSSO-r32 uses 0.453G MACs and takes 1.95 ms. Peak memory is closer to optimized MHA than the dense $O(N^2)$ formula would suggest because the benchmark uses an optimized SDPA path related to memory-efficient exact

attention kernels [7]. The practical advantage is therefore clearest in arithmetic cost and long-sequence latency, rather than peak memory at every sequence length.

5.6 Vision and Diffusion Boundary Results

Table 4 reports ViT-style [10] vision encoder results on CIFAR-100 [12] and an ImageNet subset [8]. The controlled vision runs use RandAugment [6], Mixup [28], and CutMix [26], following data-efficient vision-transformer training practice [20]. On CIFAR-100, the LSSO models trade accuracy for cost and do not match Nyströmformer. On ImageNet-100, MHA is about 1.7 points stronger in the single-seed run, whereas LSSO cuts mixer MACs by roughly three to four times. Vision therefore gives a more mixed picture than retrieval: the operator remains competitive, but its efficiency gain is not free.

Table 4: Vision encoder results. CIFAR-100 is 3 seeds; ImageNet-100 is one seed.

Dataset	Model	MACs (G)	Top-1	Top-5
CIFAR-100	MHA	0.0665	55.40 $\pm$ 0.27	—
CIFAR-100	Nyströmformer	0.0389	59.98 $\pm$ 0.58	—
CIFAR-100	LSSO-r16	0.0249	54.79 $\pm$ 0.14	—
CIFAR-100	LSSO-r32	0.0385	55.38 $\pm$ 0.18	—
ImageNet-100	MHA	0.520	82.26	95.54
ImageNet-100	LSSO-r16	0.137	80.54	94.88
ImageNet-100	LSSO-r32	0.174	80.58	94.76

We also tested a pure-LSSO latent diffusion backbone on ImageNet-100. The available experiment measures only noise-prediction validation MSE, not FID, KID, or sample quality. LSSO reaches a similar optimization regime at lower theoretical mixer cost, but this evidence is not sufficient to claim that it can replace attention in a generative model.

6 Discussion

The experiments suggest that LSSO is best viewed as a bidirectional global correction operator, not as a cheaper approximation to softmax attention. This distinction appears especially useful in retrieval, where an encoder must compress evidence from the whole sequence into a representation but does not need causal next-token routing. A small learned subspace can influence all tokens through

the solve without constructing an all-pairs value-routing matrix.

The same construction is less natural for causal decoding. Applying a causal mask breaks the simple low-rank Woodbury form, and an autoregressive extension would require prefix-restricted solves or a different cached state. The diffusion experiment points to a second boundary: global semantic mixing alone may not recover the fine-grained routing that generation demands.

7 Limitations

The theory treats U as fixed and therefore does not establish the stability of full end-to-end optimization. Rank is also a genuine capacity choice: some tasks occupy only a compact subspace, whereas SciFact uses much more of rank 32. Finally, the vision and diffusion results are narrower than the retrieval evidence. ImageNet-100 has only one seed, and the diffusion study lacks sample-quality metrics. These experiments establish useful operating points, not universal superiority over attention or other efficient mixers.

8 Conclusion

LSSO replaces pairwise value routing with a learned low-rank global solve. The positive shift makes the operator well posed, and Woodbury makes the solve practical when $r \ll N$. In our experiments this inductive bias is particularly effective for retrieval: LSSO preserves MHA-level transfer while substantially reducing mixer computation. Vision and diffusion results are less decisive, which helps delimit the present claim: LSSO is a promising efficient mixer for bidirectional encoders, rather than a general replacement for attention.

References

- [1] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. *Advances in Neural Information Processing Systems*, 2019.
- [2] Payal Bajaj, Daniel Campos, Nick Craswell, Li Deng, Jianfeng Gao, Xiaodong Liu, Rangan Majumder, Andrew McNamara, Bhaskar Mitra, Tri Nguyen, Mir Rosenberg, Xia Song, Alina Stoica, Saurabh Tiwary, and Tong Wang. MS MARCO: A human generated machine reading comprehension dataset. *arXiv preprint arXiv:1611.09268*, 2016.
- [3] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- [4] Vera Boteva, Demian Gholipour Ghalandari, Artem Sokolov, and Stefan Riezler. A full-text learning to rank dataset for medical information retrieval. In *Advances in Information Retrieval*, pages 716–722, 2016. doi: 10.1007/978-3-319-30671-1_58.
- [5] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, David Belanger, Lucy Colwell, and Adrian Weller. Rethinking attention with performers. In *International Conference on Learning Representations*, 2021.
- [6] Ekin D. Cubuk, Barret Zoph, Jonathon Shlens, and Quoc V. Le. Randaugment: Practical automated data augmentation with a reduced search space. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, pages 702–703, 2020. doi: 10.1109/CVPRW50498.2020.00359.
- [7] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35, 2022.
- [8] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 248–255, 2009. doi: 10.1109/CVPR.2009.5206848.
- [9] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, pages 4171–4186, 2019. doi: 10.18653/v1/N19-1423.
- [10] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. *International Conference on Learning Representations*, 2021.
- [11] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- [12] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [13] James Lee-Thorp, Joshua Ainslie, Ilya Eckstein, and Santiago Ontañón. Fnet: Mixing tokens with fourier transforms. In *Proceedings of NAACL-HLT*, pages 4296–4313, 2022. doi: 10.18653/v1/2022.naacl-main.319.

- [14] Hanxiao Liu, Zihang Dai, David R. So, and Quoc V. Le. Pay attention to mlps. *Advances in Neural Information Processing Systems*, 2021.
- [15] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.
- [16] Macedo Maia, Siegfried Handschuh, André Freitas, Brian Davis, Ross McDermott, Manel Zarrouk, and Alexandra Balahur. WWW’18 open challenge: Financial opinion mining and question answering. In *Companion Proceedings of The Web Conference 2018*, pages 1941–1942, 2018.
- [17] Michael Poli, Stefano Massaroli, Eric Nguyen, Daniel Y. Fu, Tri Dao, Stephen Baccus, Yoshua Bengio, Stefano Ermon, and Christopher Ré. Hyena hierarchy: Towards larger convolutional language models. *International Conference on Machine Learning*, 2023.
- [18] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 1, 2021.
- [19] Ilya O. Tolstikhin, Neil Houlsby, Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Thomas Unterthiner, Jessica Yung, Andreas Steiner, Daniel Keysers, Jakob Uszkoreit, Mario Lucic, and Alexey Dosovitskiy. Mlp-mixer: An all-mlp architecture for vision. *Advances in Neural Information Processing Systems*, 2021.
- [20] Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, and Hervé Jégou. Training data-efficient image transformers and distillation through attention. In *International Conference on Machine Learning*, pages 10347–10357, 2021.
- [21] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 2017.
- [22] David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of EMNLP*, pages 7534–7550, 2020. doi: 10.18653/v1/2020.emnlp-main.609.
- [23] Sinong Wang, Belinda Z. Li, Madian Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention with linear complexity. *arXiv preprint arXiv:2006.04768*, 2020.
- [24] Max A. Woodbury. Inverting modified matrices. Technical Report Memorandum Report 42, Statistical Research Group, Princeton University, 1950.
- [25] Yunsang Xiong, Zhanpeng Zeng, Rudras Chakraborty, Mingxing Tan, Glenn Fung, Yin Li, and Vikas Singh. Nyströmformer: A nyström-based algorithm for approximating self-attention. In *AAAI Conference on Artificial Intelligence*, 2021. doi: 10.1609/aaai.v35i16.17664.
- [26] Sangdoo Yun, Dongyoon Han, Seong Joon Oh, Sanghyuk Chun, Junsuk Choe, and Youngjoon Yoo. Cutmix: Regularization strategy to train strong classifiers with localizable features. In *IEEE/CVF International Conference on Computer Vision*, pages 6023–6032, 2019. doi: 10.1109/ICCV.2019.00612.
- [27] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alben, Santiago Ontiveros, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big bird: Transformers for longer sequences. *Advances in Neural Information Processing Systems*, 2020.
- [28] Hongyi Zhang, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations*, 2018.

A Full Woodbury Derivation

Using the Woodbury identity with $A = \mu I_N$, $C = \gamma I_r$, and $V = U^\top$ gives

$$(\mu I + \gamma U U^\top)^{-1} = \mu^{-1} I - \mu^{-1} U (\gamma^{-1} I + \mu^{-1} U^\top U)^{-1} U^\top \mu^{-1}.$$

Multiplying the small inverse by γ avoids an explicit γ^{-1} and gives the implementation form

$$Y = \mu^{-1} C - \mu^{-1} U (I + \gamma \mu^{-1} U^\top U)^{-1} (\gamma \mu^{-1} U^\top C),$$

which remains valid at $\gamma = 0$.