

LSSO: Learnable Low-Rank Sylvester Solves for Efficient Bidirectional Encoders

Zhaoyang Wang

zhaoy@uir.edu.cn

<https://github.com/Yang916-yy/LSSO>

Abstract

We introduce LSSO, a token mixer for bidirectional encoders built around a learned low-rank linear solve. Rather than routing values through an explicit pairwise attention matrix, LSSO constructs an input-dependent relation field and solves

$$(\mu I + \gamma U(X)U(X)^\top)Y = C(X).$$

This is the left-sided case of a Sylvester equation. The positive shift makes the system well posed, while the Woodbury identity reduces the computation to an $r \times r$ solve. For fixed relation features, the operator has a bounded spectral response and minimizes a strongly convex correction energy. Across retrieval, MS MARCO-to-BEIR transfer, and vision classification, LSSO matches or approaches attention-based encoders while using substantially fewer mixer MACs. Ablations, rank pruning, and sequence-scaling experiments further show when the global correction is useful and where the current formulation falls short. We also derive an exact causal prefix form: causal masking can be rewritten as prefix low-rank statistics and batched small solves, with a fixed-size solve-state cache for decoding.

1 Introduction

Self-attention [22] is the standard token mixer in modern bidirectional encoders. Given query, key, and value projections, it computes

$$Y_{\text{attn}} = \text{softmax}(QK^\top/\sqrt{d})V,$$

and lets every token route information from every other token. That flexibility comes with quadratic score and value-routing costs. Most efficient attention methods reduce the cost by approximating or sparsifying the routing matrix. LSSO asks a different question: can a global token interaction be learned without constructing such a matrix at all?

For an input sequence $X \in \mathbb{R}^{B \times N \times D}$, LSSO learns low-rank relation features $U(X)$ and a content state $C(X)$. Each head solves

$$(\mu I_N + \gamma UU^\top)Y = C, \quad \mu > 0, \quad \gamma \geq 0,$$

and projects the result back to the model dimension. Although $UU^\top$ is positive semidefinite, its entries $u_i^\top u_j$ may be positive or negative. The system can therefore encode signed token relations while remaining invertible by construction.

The guiding interpretation is:

LSSO = isotropic content pass-through
— low-rank global correction.

Attention learns a row-normalized routing rule; LSSO learns the relation features of a positive-definite inverse filter. The two mixers therefore have different inductive biases even when their input and output shapes match.

The main contributions are:

1. We formulate bidirectional token mixing as a learned, positive-shifted low-rank solve whose implementation only solves an $r \times r$ system.
2. We characterize the fixed- U operator through its spectrum, conditioning, norm bound, and equivalent convex energy.
3. We derive a causal prefix-LSSO form that avoids an explicit triangular attention matrix by using prefix low-rank statistics and fixed-size solve-state caches.
4. We evaluate the mixer on retrieval, zero-shot transfer, vision, rank pruning, diagnostics, and long-sequence scaling, including settings where LSSO does not outperform attention.

2 Related Work

Efficient attention. Linformer [24], Performer [5], Nyströmformer [26], Longformer [3], and BigBird [28] reduce the cost of self-attention through low-rank projections, kernel approximations, landmark approximations, or sparse attention patterns. Despite these differences, they retain attention as the underlying interaction rule. LSSO instead replaces that rule with an input-dependent inverse filter.

Token mixers beyond attention. MLP-Mixer [20], gMLP [14], FNet [13], Hyena [17], and state-space models such as Mamba [11] demonstrate that useful token mixing does not require dense attention. LSSO explores a complementary construction: an implicit low-rank solve designed specifically for non-causal encoders.

Implicit layers and learned solves. Deep equilibrium models [1], implicit layers, graph smoothing, and regularized least-squares methods place linear or nonlinear solves inside neural networks. LSSO is narrower than a general equilibrium layer: it applies one input-conditioned linear solve with a deliberately constrained positive-definite matrix.

3 Method

3.1 Encoder Block

LSSO is used in the token-mixer slot of an encoder block. In a pre-normalized Transformer-style block this is simply

$$X' = X + \text{LSSO}(\text{LN}(X)), \quad X_{\text{out}} = X' + \text{MLP}(\text{LN}(X')).$$

The replacement is local: its input and output shapes match MHA, so the residual path, normalization, positional encoding, feed-forward block, and task head can be kept. In the BERT-style retrieval experiments we preserve the post-LN BERT block used by the baseline; in the ViT-style vision experiments we use the pre-LN form above. The comparison therefore changes the mixer without changing the surrounding block convention.

3.2 LSSO Definition

For each head, let

$$U = XW_u \in \mathbb{R}^{B \times N \times r}, \quad C = XW_c \in \mathbb{R}^{B \times N \times d_h}.$$

Each head has its own pair of scalar solve parameters. We RMS-normalize U along its rank dimension by default,

$$U \leftarrow U / \sqrt{\text{mean}_r(U^2) + \epsilon},$$

and solve

$$(\mu I_N + \gamma U U^\top) Y = C, \quad (1)$$

where

$$\mu = \text{softplus}(\theta_\mu) + \epsilon, \quad \gamma = \gamma_{\max} \sigma(\theta_\gamma).$$

The default experiments use $\gamma_{\max} = 0.3$ and $\theta_\gamma = -4$, which gives an initial γ/μ of roughly 7.8×10^{-3} when $\theta_\mu = 0$. More conservative settings such as $\gamma_{\max} = 0.1$, $\theta_\gamma = -6$ make the global correction nearly vanish at initialization; in preliminary runs they consistently underused the global term.

The head outputs are concatenated and projected:

$$\text{LSSO}(X) = YW_o.$$

For padded sequences, invalid tokens are masked by zeroing the corresponding rows of both U and C before the solve and by zeroing the output at masked positions. This mirrors key-padding behavior in bidirectional attention while preserving the low-rank Woodbury structure over the valid tokens.

3.3 Woodbury Implementation

The $N \times N$ matrix in Eq. (1) never needs to be formed or inverted. Applying the Woodbury identity [25] gives

$$Y = \mu^{-1} C - \mu^{-1} U (I_r + \gamma \mu^{-1} U^\top U)^{-1} (\gamma \mu^{-1} U^\top C). \quad (2)$$

This expression remains valid at $\gamma = 0$ and requires only an $r \times r$ solve. Appendix A gives the full derivation. The released implementation fuses the U and C projections into one linear map; this is algebraically identical to using separate W_u and W_c . The small positive-definite system is solved by Cholesky in float32 under fp16/bf16 autocast, then cast back to the activation dtype. The package also provides a custom autograd path whose backward pass reuses the same Woodbury solve on the adjoint, avoiding a dense $N \times N$ inverse in both directions. The core solve terms cost

$$O(Nr^2 + Nr d_h + r^2 d_h + r^3)$$

per head, where $r^2 d_h$ accounts for applying the Cholesky factors to the d_h right-hand sides. Including projections, the mixer-only MAC estimate used in our tables is

$$ND(Hr + D) + ND^2 + HN r^2 + 2NrD + Dr^2 + Hr^3.$$

Here the final term conservatively counts each small factorization as r^3 MAC-equivalents. For MHA we count QKV projections, attention scores, attention-value product, and output projection:

$$4ND^2 + 2N^2 D.$$

4 Theoretical Properties

We first study the linear map induced by a fixed U ; the input-dependent operator is evaluated empirically in Section 6.

Proposition 1 (Well-posedness). *Let $A = \mu I_N + \gamma U U^\top$ with $\mu > 0$ and $\gamma \geq 0$. Then A is symmetric positive definite. Therefore, for any C , the system $AY = C$ has a unique solution.*

Proof. For any nonzero z ,

$$z^\top A z = \mu \|z\|_2^2 + \gamma \|U^\top z\|_2^2 \geq \mu \|z\|_2^2 > 0.$$

Proposition 2 (Spectral shrinkage). *Let $UU^\top = Q\Lambda Q^\top$ be an eigendecomposition over the nonzero eigen-directions of $UU^\top$, with $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_k)$ and $\lambda_i > 0$. Let $P_\perp = I - QQ^\top$. Then*

$$A^{-1} = Q \text{diag}\left(\frac{1}{\mu + \gamma\lambda_i}\right) Q^\top + \mu^{-1}P_\perp.$$

Thus LSSO applies an inverse spectral filter on the span of U and an isotropic scale on its orthogonal complement. This gives a more precise view of the distinction from row-normalized attention.

Proposition 3 (Bounded fixed- U operator norm). *For fixed U ,*

$$\|Y\|_F \leq \mu^{-1}\|C\|_F.$$

For a multi-head layer, let Y_{cat} and C_{cat} denote the concatenated head outputs and inputs, and let $\mu_{\min} = \min_h \mu_h$. With the output projection applied on the right,

$$\|Y_{\text{cat}}W_o\|_F \leq \mu_{\min}^{-1}\|C_{\text{cat}}\|_F\|W_o\|_2.$$

Proof. Since $A \succeq \mu I$, $\lambda_{\min}(A) \geq \mu$ and $\|A^{-1}\|_2 \leq \mu^{-1}$. The Frobenius bound follows from $Y = A^{-1}C$. Summing the squared per-head bounds gives $\|Y_{\text{cat}}\|_F \leq \mu_{\min}^{-1}\|C_{\text{cat}}\|_F$, and $\|Y_{\text{cat}}W_o\|_F \leq \|Y_{\text{cat}}\|_F\|W_o\|_2$ completes the result.

Proposition 4 (Conditioning control). *For $A = \mu I + \gamma UU^\top$,*

$$\kappa_2(A) \leq 1 + \frac{\gamma}{\mu}\|U\|_2^2.$$

Propositions 3 and 4 motivate the diagnostics used in our experiments. The ratio γ/μ directly scales the possible condition-number increase; effective rank measures how much of the rank subspace is used; and correction ratio measures how strongly the low-rank correction participates in the output. RMS normalization of U helps control the scale of U , although it is not by itself a dimension-free spectral-norm guarantee.

Theorem 1 (Energy minimization equivalence). *For $\mu > 0$ and $\gamma \geq 0$, the LSSO output is the unique minimizer of*

$$\mathcal{E}(Y) = \frac{\mu}{2}\|Y\|_F^2 + \frac{\gamma}{2}\|U^\top Y\|_F^2 - \langle C, Y \rangle.$$

Proof. The Hessian of $\mathcal{E}$ is $\mu I + \gamma UU^\top$, which is positive definite by Proposition 1. Setting the gradient to zero gives $(\mu I + \gamma UU^\top)Y = C$.

Equivalently, up to an additive constant,

$$\mathcal{E}(Y) = \frac{\mu}{2}\|Y - \mu^{-1}C\|_F^2 + \frac{\gamma}{2}\|U^\top Y\|_F^2.$$

The content projection supplies the local state $\mu^{-1}C$, and the second term adjusts that state using sequence-wide information.

Proposition 5 (No-global degeneration). *If $\gamma = 0$, then*

$$Y = \mu^{-1}C.$$

No token mixing occurs except for pointwise projections.

5 Causal Prefix-LSSO

The LSSO layer studied above is bidirectional: all outputs come from one full-sequence solve. A causal variant should not be obtained by inserting an $N \times N$ triangular mask into Eq. (1); that would obscure the low-rank structure. Instead, define each output by the solve on its own prefix. For position t ,

$$y_t^\top = [(\mu I_t + \gamma U_{1:t}U_{1:t}^\top)^{-1}C_{1:t}]_{t,:}. \quad (3)$$

This is an exact causal operator: y_t depends only on tokens $1, \dots, t$. The inclusive prefix in Eq. (3) matches standard causal attention, where a token can read itself. An exclusive version is obtained by using statistics from $1, \dots, t-1$.

For one head, write the rows of U and C as $u_t^\top \in \mathbb{R}^{1 \times r}$ and $c_t^\top \in \mathbb{R}^{1 \times d_h}$. Let $U_{1:t}$ and $C_{1:t}$ contain the first t rows, and set

$$\alpha = \frac{\gamma}{\mu}.$$

Define the prefix low-rank statistics

$$G_t = \sum_{j \leq t} u_j u_j^\top = U_{1:t}^\top U_{1:t} \in \mathbb{R}^{r \times r}, \quad (4)$$

$$M_t = \sum_{j \leq t} u_j c_j^\top = U_{1:t}^\top C_{1:t} \in \mathbb{R}^{r \times d_h}. \quad (5)$$

Applying Eq. (2) to each prefix gives

$$y_t^\top = \mu^{-1} [c_t^\top - \alpha u_t^\top (I_r + \alpha G_t)^{-1} M_t]. \quad (6)$$

Thus causal training or prefill can be written as a parallel prefix scan: compute the per-token outer products $u_t u_t^\top$ and $u_t c_t^\top$, scan them over the sequence to obtain G_t and M_t for every position, and solve the batched $r \times r$ systems in Eq. (6). The only sequential primitive is the scan; the solves are independent across positions once the prefix statistics are available.

For autoregressive decoding, one may simply cache G_t and M_t :

$$G_t = G_{t-1} + u_t u_t^\top, \quad (7)$$

$$M_t = M_{t-1} + u_t c_t^\top. \quad (8)$$

The cache size per layer and head is $r^2 + r d_h$, independent of the generated length. The next token can then be read with Eq. (6). This is analogous in role to a KV cache, but it stores solve statistics rather than past keys and values.

5.1 Rank-One Recurrent Readout

The prefix solve also admits an equivalent recurrence that updates the inverse small system directly. Let

$$S_t = I_r + \alpha G_t, \quad T_t = M_t.$$

The two recurrent states are

$$P_t = S_t^{-1} \in \mathbb{R}^{r \times r}, \quad Z_t = P_t T_t \in \mathbb{R}^{r \times d_h}.$$

Proposition 6 (Exact prefix recurrence). *Initialize $P_0 = I_r$ and $Z_0 = 0$. For token t , define*

$$q_t = P_{t-1}u_t, \quad \delta_t = 1 + \alpha u_t^\top q_t.$$

Then the updates

$$P_t = P_{t-1} - \frac{\alpha}{\delta_t} q_t q_t^\top, \quad (9)$$

$$Z_t = Z_{t-1} + \frac{q_t}{\delta_t} (c_t^\top - \alpha u_t^\top Z_{t-1}) \quad (10)$$

are well defined, and the output at step t

$$y_t^\top = \mu^{-1} (c_t^\top - \alpha u_t^\top Z_t) \quad (11)$$

is exactly the last row of the prefix solve

$$(\mu I_t + \gamma U_{1:t} U_{1:t}^\top) Y_{1:t} = C_{1:t}.$$

Proof. Since S_{t-1} is positive definite, so is P_{t-1} . Therefore $\delta_t = 1 + \alpha u_t^\top P_{t-1} u_t \geq 1$, and the denominator is positive. Moreover,

$$S_t = S_{t-1} + \alpha u_t u_t^\top.$$

Applying the Sherman–Morrison rank-one inverse identity [18] gives Eq. (9), hence $P_t = S_t^{-1}$. It also gives

$$P_t u_t = \frac{q_t}{\delta_t}.$$

Using $T_t = T_{t-1} + u_t c_t^\top$ and $Z_{t-1} = P_{t-1} T_{t-1}$,

$$\begin{aligned} Z_t &= P_t (T_{t-1} + u_t c_t^\top) \\ &= Z_{t-1} - \frac{\alpha q_t u_t^\top Z_{t-1}}{\delta_t} + \frac{q_t c_t^\top}{\delta_t}, \end{aligned}$$

which is Eq. (10). Finally, applying the Woodbury expression in Eq. (2) to the length- t prefix yields

$$Y_{1:t} = \mu^{-1} (C_{1:t} - \alpha U_{1:t} S_t^{-1} T_t).$$

Taking its last row and substituting $Z_t = S_t^{-1} T_t$ proves Eq. (11).

Per head, the recurrence stores $r^2 + r d_h$ values, independent of prefix length, and costs $O(r^2 + r d_h)$ per token beyond the projections. In this operational sense it resembles a state-space sequence model: each token updates a fixed-size state and produces an output without retaining the prefix. It is not a linear state-space recurrence. The transitions in Eqs. (9)–(10) are input-dependent rational maps, similar in that respect to selective state-space models [11].

Equation (11) gives the last-row solution for the prefix available at step t . It does not recover the t -th row of a longer bidirectional solve, since appending a token also changes earlier solved states. The recurrence therefore implements a causal *prefix LSSO*; the parallel Woodbury form computes the bidirectional LSSO studied in this paper. The algebra therefore suggests three compatible implementations of the same low-rank construction: parallel full-context encoding, prefix-scan causal training, and fixed-state recurrent decoding.

6 Experiments

6.1 Setup

All main retrieval and vision comparisons train randomly initialized encoders. The retrieval backbone follows a BERT-style bidirectional encoder layout [9]. We reuse standard tokenization but train every encoder weight from scratch. Models are optimized with AdamW [15]. LSSO uses $\gamma_{\max} = 0.3$ and $\theta_\gamma = -4$ unless stated otherwise. Because the feed-forward network is identical across models, tables report mixer-only MACs and isolate the cost of the component being replaced.

Retrieval protocol. The dual encoder has $D = 256$, depth 8, 8 heads, and dropout 0.1. Queries and documents are truncated to 64 and 512 tokens, respectively, and representations are obtained by masked mean pooling. We train with in-batch contrastive cross-entropy at temperature 0.05, using batch size 64, AdamW with learning rate 3×10^{-4} and weight decay 0.01, 5% linear warmup followed by cosine decay, and gradient clipping at 1.0. Each target-dataset run uses at most 50,000 training pairs for 20 epochs; the checkpoint with the highest validation R@10 is evaluated. All models use the same sampled pairs, tokenizer, pooling rule, and three random seeds.

6.2 Retrieval Main Results

Table 1 reports three-seed retrieval results on FIQA [16], NFCorpus [4], and SciFact [23], using their BEIR retrieval formulations [19]. The baseline is a BERT-style MHA encoder, and Nyströmformer uses the authors’ attention module. LSSO-r16 and LSSO-r32 use the same model dimension, depth, number of heads, optimizer, and training schedule.

Table 1: Retrieval main results. Models are randomly initialized BERT-style encoders with $D = 256$, depth 8, heads 8, and document length 512. MACs are mixer-only document-side MACs.

Dataset	Model	Params (M)	Mixer MACs (G)	R@10	MRR@10
FIQA	MHA	14.33	2.147	0.2145 $\pm$ 0.0135	0.0987 $\pm$ 0.0032
FIQA	Nyströmformer	14.33	1.301	0.2428 $\pm$ 0.0039	0.1190 $\pm$ 0.0071
FIQA	LSSO-r16	13.54	0.714	0.2258 $\pm$ 0.0125	0.1126 $\pm$ 0.0061
FIQA	LSSO-r32	13.80	0.910	0.2387 $\pm$ 0.0183	0.1150 $\pm$ 0.0035
NFCorpus	MHA	14.33	2.147	0.5439 $\pm$ 0.0187	0.3809 $\pm$ 0.0145
NFCorpus	Nyströmformer	14.33	1.301	0.5841 $\pm$ 0.0125	0.4323 $\pm$ 0.0132
NFCorpus	LSSO-r16	13.54	0.714	0.5686 $\pm$ 0.0125	0.4121 $\pm$ 0.0036
NFCorpus	LSSO-r32	13.80	0.910	0.5841 $\pm$ 0.0036	0.4333 $\pm$ 0.0220
SciFact	MHA	14.33	2.147	0.6789 $\pm$ 0.0190	0.5994 $\pm$ 0.0138
SciFact	Nyströmformer	14.33	1.301	0.7267 $\pm$ 0.0133	0.6313 $\pm$ 0.0033
SciFact	LSSO-r16	13.54	0.714	0.7022 $\pm$ 0.0107	0.6214 $\pm$ 0.0119
SciFact	LSSO-r32	13.80	0.910	0.7089 $\pm$ 0.0117	0.6247 $\pm$ 0.0080

Relative to MHA, LSSO-r16 removes 66.8% of mixer MACs and LSSO-r32 removes 57.6%. Both improve the mean MHA result on every dataset in the table. Nyströmformer remains the strongest baseline on several metrics, although in our implementation its lower theoretical MAC count did not translate into higher measured throughput.

6.3 MS MARCO to BEIR Transfer

We next ask whether the retrieval result survives a change of domain. Each model is pretrained once on MS MARCO [2] and evaluated without target-task finetuning on five BEIR datasets [19] covering finance, biomedicine, scientific fact checking, argument retrieval, and COVID literature. Unlike the preceding experiments, this setting does not allow the encoder to adapt to each evaluation collection.

All four models use the same BERT-style architecture, initialization protocol, MS MARCO training data, tokenizer, pooling rule, optimizer, three seeds, and document length 512; only the token mixer changes. For each seed, we average each metric across the five target datasets, then report the mean and standard deviation of these macro averages. Since total parameter counts are similar (13.53–14.33M), Table 2 emphasizes mixer MACs. The released artifacts contain the corresponding per-dataset and per-seed values. Each model is trained for two epochs on 500,000 MS MARCO pairs with batch size 192, bf16 autocast, learning rate 3×10^{-4} , 5% warmup, and the same temperature and clipping settings as the target-dataset runs. Corpus embeddings are computed once per checkpoint and searched by inner product; no labels or optimizer steps from the five BEIR targets are used.

Table 2: Zero-shot BEIR transfer after MS MARCO pre-training. Retrieval metrics are five-dataset macro averages, reported as mean $\pm$ standard deviation over three seeds. MACs are mixer-only at document length 512. Best values are bold.

Model	Mixer MACs (G)	nDCG@10	Recall@100	MRR@10
MHA	2.147	0.22945 $\pm$ 0.00208	0.36061 $\pm$ 0.00365	0.30225$\pm$0.00442
Nyströmformer	1.270 (−40.9%)	0.20523 $\pm$ 0.00495	0.33456 $\pm$ 0.00837	0.27115 $\pm$ 0.01046
LSSO-r16	0.714 (−66.8%)	0.22973$\pm$0.00232	0.35889 $\pm$ 0.00226	0.29382 $\pm$ 0.01446
LSSO-r32	0.910 (−57.6%)	0.22940 $\pm$ 0.00337	0.36367$\pm$0.00830	0.29858 $\pm$ 0.00624

Both LSSO variants preserve MHA-level transfer at a much lower mixer cost. LSSO-r16 slightly exceeds the MHA macro nDCG@10 (0.22973 vs. 0.22945) while using 66.8% fewer mixer MACs. LSSO-r32 has the strongest macro Recall@100 (0.36367) while using 57.6% fewer mixer MACs. Nyströmformer is also cheaper than MHA, but trails both MHA and LSSO on all three macro metrics.

No single rank wins every metric. LSSO-r16 offers the best nDCG@10–cost trade-off, LSSO-r32 reaches the highest Recall@100, and MHA retains the best MRR@10. Taken together, the two LSSO settings lie on the same transfer frontier as MHA while requiring considerably less mixer computation.

6.4 Ablations and Diagnostics

Table 3 aligns the empirical ablations with the theory. The no-global variant sets $\gamma = 0$, reducing the operator to $\mu^{-1}C$ and removing token mixing as stated in Proposition 5. The fixed-scale variant freezes both θ_μ and θ_γ at initialization, while the no-normalization variant removes the row-wise RMS normalization of U . The r8 and r4 rows

are independently trained full LSSO models at smaller rank, rather than pruned r32 checkpoints. Removing RMS normalization weakens the global correction, especially on SciFact, while the smaller-rank runs expose the capacity needed by the learned relation subspace.

Table 3: Retrieval ablations. Values are mean $\pm$ standard deviation over three seeds.

Dataset	Variant	R@10	MRR@10	γ/μ	Correction ratio
FIQA	no-global r32	0.2052 $\pm$ 0.0137	0.0911 $\pm$ 0.0019	0.0000	0.0000
FIQA	fixed μ/γ r32	0.2428 $\pm$ 0.0126	0.1136 $\pm$ 0.0028	0.0078	0.2130
FIQA	no U RMS norm r32	0.2160 $\pm$ 0.0101	0.1007 $\pm$ 0.0008	0.0085	0.2050
FIQA	full r8	0.2088 $\pm$ 0.0208	0.0923 $\pm$ 0.0139	0.0100	0.1059
FIQA	full r4	0.2042 $\pm$ 0.0045	0.0979 $\pm$ 0.0036	0.0109	0.0742
SciFact	no-global r32	0.6767 $\pm$ 0.0120	0.5811 $\pm$ 0.0058	0.0000	0.0000
SciFact	fixed μ/γ r32	0.7089 $\pm$ 0.0102	0.6088 $\pm$ 0.0066	0.0078	0.2631
SciFact	no U RMS norm r32	0.6933 $\pm$ 0.0133	0.5972 $\pm$ 0.0150	0.0078	0.0471
SciFact	full r8	0.7067 $\pm$ 0.0173	0.6118 $\pm$ 0.0064	0.0078	0.1149
SciFact	full r4	0.7022 $\pm$ 0.0069	0.6146 $\pm$ 0.0107	0.0078	0.0766

Figure 1 shows the same quantities layer by layer. The correction ratio remains nonzero on every task, so the trained models do not collapse to the local projection. Effective rank varies more substantially: SciFact uses a larger fraction of the rank-32 subspace than FIQA, NFCorpus, or CIFAR-100.




Figure 1: Layer-wise LSSO diagnostics from trained LSSO-r32 checkpoints.

6.5 Inference-Time Rank Pruning

The low effective ranks in Figure 1 suggest that a trained rank-32 model may not need every coordinate at inference. For each input and head, we rank coordinates by mean squared activation in U and retain the top k before the Woodbury solve. Table 4 evaluates this procedure directly on the three-seed LSSO-r32 retrieval checkpoints, without finetuning after pruning.

Table 4: Inference-time pruning of trained LSSO-r32 models. “Save” is the theoretical mixer-MAC reduction of a compact rank- k implementation relative to rank 32; the dynamic evaluator is used only to measure quality.

Keep rank	Save	FIQA R@10	NFCorpus R@10	SciFact R@10
32	0.0%	0.2387 $\pm$ 0.0183	0.5841 $\pm$ 0.0036	0.7089 $\pm$ 0.0117
16	21.6%	0.2377 $\pm$ 0.0107	0.5831 $\pm$ 0.0047	0.7056 $\pm$ 0.0135
8	31.5%	0.2361 $\pm$ 0.0149	0.5810 $\pm$ 0.0078	0.6922 $\pm$ 0.0139
4	36.3%	0.2269 $\pm$ 0.0157	0.5759 $\pm$ 0.0054	0.6856 $\pm$ 0.0139

Keeping 16 of 32 coordinates changes R@10 by less than 0.004 on every dataset while reducing the theoretical compact mixer cost by 21.6%. More aggressive pruning remains mild on FIQA and NFCorpus but degrades SciFact more clearly, consistent with its higher measured effective rank. These savings require exporting or executing the reduced-rank operator; merely masking coordinates inside the original rank-32 tensors does not reduce projection cost.

6.6 Sequence Scaling

We benchmark a single mixer layer at sequence lengths from 128 to 2048, using batch size 8, $D = 256$, 8 heads, bf16 autocast, and no optimizer update. Figure 2 reports forward latency, forward-plus-backward latency, incremental peak memory, and theoretical mixer MACs.




Figure 2: Sequence scaling benchmark on an NVIDIA RTX 5070 Ti. MACs are theoretical per-sample single-layer mixer MACs; latency and memory use batch 8.

At $N = 2048$, LSSO-r16 uses 0.357G mixer MACs compared with 2.684G for MHA, an 86.7% reduction, and reduces forward+backward latency from 2.61 ms to 1.88 ms. LSSO-r32 uses 0.454G MACs and takes 1.95 ms. Peak memory is closer to optimized MHA than the dense $O(N^2)$ formula would suggest because the benchmark uses an optimized SDPA path related to memory-efficient exact attention kernels [7]. The practical advantage is therefore clearest in arithmetic cost and long-sequence latency, rather than peak memory at every sequence length.

6.7 Vision and Diffusion Boundary Results

Table 5 reports ViT-style [10] vision encoder results on CIFAR-100 [12] and an ImageNet subset [8]. The controlled vision runs use RandAugment [6], Mixup [29], and CutMix [27], following data-efficient vision-transformer training practice [21]. CIFAR-100 uses 32×32 images, patch size 2, $D = 96$, depth 3, 6 heads, batch size 256, and 100 epochs. Its augmentation consists of random crop and horizontal flip, RandAugment, Mixup 0.2, CutMix 0.5, and label smoothing 0.1. ImageNet-100 uses 224×224 images, patch size 8, $D = 256$, depth 8, 8 heads, batch size 256, and 120 epochs, with Mixup 0.8, CutMix 1.0, RandAugment, and label smoothing 0.1. Both use AdamW, cosine decay, and identical settings across mixers within a dataset.

On CIFAR-100, the LSSO models trade accuracy for cost and do not match Nyströmformer. On ImageNet-100, MHA is about 1.7 points stronger in the single-seed

run, whereas LSSO cuts mixer MACs by roughly three to four times. Vision therefore gives a more mixed picture than retrieval: the operator remains competitive, but its efficiency gain is not free.

Table 5: Vision encoder results. MACs cover all mixer layers in the encoder. CIFAR-100 results use three seeds; ImageNet-100 uses one.

Dataset	Model	MACs (G)	Top-1	Top-5
CIFAR-100	MHA	0.0665	55.40 ± 0.27	—
CIFAR-100	Nyströmformer	0.0389	59.98 ± 0.58	—
CIFAR-100	LSSO-r16	0.0250	54.79 ± 0.14	—
CIFAR-100	LSSO-r32	0.0388	55.38 ± 0.18	—
ImageNet-100	MHA	4.162	82.26	95.54
ImageNet-100	LSSO-r16	1.093	80.54	94.88
ImageNet-100	LSSO-r32	1.391	80.58	94.76

We also tested a pure-LSSO latent diffusion backbone on ImageNet-100. This experiment reports noise-prediction validation MSE rather than FID, KID, or sample quality. LSSO reaches a similar optimization regime at lower theoretical mixer cost, but this evidence is not sufficient to claim that it can replace attention in a generative model.

7 Discussion

The experiments suggest that LSSO is best viewed as a bidirectional global correction operator, not as a cheaper approximation to softmax attention. This distinction appears especially useful in retrieval, where an encoder must compress evidence from the whole sequence into a representation but does not need causal next-token routing. A small learned subspace can influence all tokens through the solve without constructing an all-pairs value-routing matrix. This interpretation is supported by the retrieval results in Table 1, the zero-shot transfer results in Table 2, and the nonzero correction ratios in Figure 1.

Causal prefix form. Section 5 shows that the same positive-shifted low-rank solve has an exact causal prefix form. During training or prefill, the masked dependency can be expressed as prefix statistics followed by batched small solves; during decoding, the same operator admits a fixed-size solve-state cache. This makes causal LSSO a concrete extension rather than a triangular masking trick. At the same time, the empirical evidence in this paper remains focused on bidirectional encoders. Competitive causal language modeling would still require dedicated kernels, careful low-precision state updates, and a separate study of optimization and generation quality.

8 Limitations

Rank is a genuine capacity choice: some tasks occupy only a compact subspace, whereas SciFact uses much more of rank 32 (Figure 1 and Table 4). The vision and diffusion results are also narrower than the retrieval evidence.

ImageNet-100 has only one seed, and the diffusion study lacks sample-quality metrics. These experiments establish useful operating points, not universal superiority over attention or other efficient mixers.

9 Conclusion

LSSO replaces pairwise value routing with a learned low-rank global solve. The positive shift makes the operator well posed, and Woodbury makes the solve practical when $r \ll N$ (Eq. (2) and Appendix A). In our experiments this inductive bias is particularly effective for retrieval: LSSO preserves MHA-level transfer while substantially reducing mixer computation (Tables 1 and 2). Vision and diffusion results are less decisive, which helps delimit the present claim: LSSO is a promising efficient mixer for bidirectional encoders, rather than a general replacement for attention. The causal prefix form derived in Section 5 gives a concrete route toward future streaming and autoregressive variants, but those variants require their own implementation and empirical evaluation.

References

- [1] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. *Advances in Neural Information Processing Systems*, 2019.
- [2] Payal Bajaj, Daniel Campos, Nick Craswell, Li Deng, Jianfeng Gao, Xiaodong Liu, Rangan Majumder, Andrew McNamara, Bhaskar Mitra, Tri Nguyen, Mir Rosenberg, Xia Song, Alina Stoica, Saurabh Tiwary, and Tong Wang. MS MARCO: A human generated machine reading comprehension dataset. *arXiv preprint arXiv:1611.09268*, 2016.
- [3] Iz Beltagy, Matthew E. Peters, and Arman Cohan. Longformer: The long-document transformer. *arXiv preprint arXiv:2004.05150*, 2020.
- [4] Vera Boteva, Demian Gholipour Ghalandari, Artem Sokolov, and Stefan Riezler. A full-text learning to rank dataset for medical information retrieval. In *Advances in Information Retrieval*, pages 716–722, 2016. doi: 10.1007/978-3-319-30671-1_58.
- [5] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, David Belanger, Lucy Colwell, and Adrian Weller. Rethinking attention with performers. In *International Conference on Learning Representations*, 2021.
- [6] Ekin D. Cubuk, Barret Zoph, Jonathon Shlens, and Quoc V. Le. Randaugment: Practical automated data augmentation with a reduced search space. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, pages 702–703, 2020. doi: 10.1109/CVPRW50498.2020.00359.
- [7] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. Flashattention: Fast and memory-efficient exact attention with io-awareness. *Advances in Neural Information Processing Systems*, 35, 2022.
- [8] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition*, pages 248–255, 2009. doi: 10.1109/CVPR.2009.5206848.
- [9] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, pages 4171–4186, 2019. doi: 10.18653/v1/N19-1423.
- [10] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold, Sylvain Gelly, Jakob Uszkoreit, and Neil Houlsby. An image is worth 16x16 words: Transformers for image recognition at scale. *International Conference on Learning Representations*, 2021.
- [11] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. In *First Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=tEYskw1VY2>.
- [12] Alex Krizhevsky. Learning multiple layers of features from tiny images. Technical report, University of Toronto, 2009.
- [13] James Lee-Thorp, Joshua Ainslie, Ilya Eckstein, and Santiago Ontañón. Fnet: Mixing tokens with fourier transforms. In *Proceedings of NAACL-HLT*, pages 4296–4313, 2022. doi: 10.18653/v1/2022.naacl-main.319.
- [14] Hanxiao Liu, Zihang Dai, David R. So, and Quoc V. Le. Pay attention to mlps. *Advances in Neural Information Processing Systems*, 2021.
- [15] Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.
- [16] Macedo Maia, Siegfried Handschuh, André Freitas, Brian Davis, Ross McDermott, Manel Zarrouk, and Alexandra Balahur. WWW’18 open challenge: Financial opinion mining and question answering. In *Companion Proceedings of The Web Conference 2018*, pages 1941–1942, 2018. doi: 10.1145/3184558.3192301.

- [17] Michael Poli, Stefano Massaroli, Eric Nguyen, Daniel Y. Fu, Tri Dao, Stephen Baccus, Yoshua Bengio, Stefano Ermon, and Christopher Ré. Hyena hierarchy: Towards larger convolutional language models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202, pages 28043–28078. PMLR, 2023. URL <https://proceedings.mlr.press/v202/poli23a.html>.
- [18] Jack Sherman and Winifred J. Morrison. Adjustment of an inverse matrix corresponding to a change in one element of a given matrix. *The Annals of Mathematical Statistics*, 21(1):124–127, 1950. doi: 10.1214/aoms/1177729893.
- [19] Nandan Thakur, Nils Reimers, Andreas Rücklé, Abhishek Srivastava, and Iryna Gurevych. Beir: A heterogeneous benchmark for zero-shot evaluation of information retrieval models. *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*, 1, 2021.
- [20] Ilya O. Tolstikhin, Neil Houlsby, Alexander Kolesnikov, Lucas Beyer, Xiaohua Zhai, Thomas Unterthiner, Jessica Yung, Andreas Steiner, Daniel Keysers, Jakob Uszkoreit, Mario Lucic, and Alexey Dosovitskiy. Mlp-mixer: An all-mlp architecture for vision. *Advances in Neural Information Processing Systems*, 2021.
- [21] Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, and Hervé Jégou. Training data-efficient image transformers and distillation through attention. In *International Conference on Machine Learning*, pages 10347–10357, 2021.
- [22] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 2017.
- [23] David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hananeh Hajishirzi. Fact or fiction: Verifying scientific claims. In *Proceedings of EMNLP*, pages 7534–7550, 2020. doi: 10.18653/v1/2020.emnlp-main.609.
- [24] Sinong Wang, Belinda Z. Li, Madian Khabsa, Han Fang, and Hao Ma. Linformer: Self-attention with linear complexity. *arXiv preprint arXiv:2006.04768*, 2020.
- [25] Max A. Woodbury. Inverting modified matrices. Technical Report Memorandum Report 42, Statistical Research Group, Princeton University, 1950.
- [26] Yunyang Xiong, Zhanpeng Zeng, Rudrasis Chakraborty, Mingxing Tan, Glenn Fung, Yin Li, and Vikas Singh. Nyströmformer: A nyström-based algorithm for approximating self-attention. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 14138–14148, 2021. doi: 10.1609/aaai.v35i16.17664.
- [27] Sangdoo Yun, Dongyoon Han, Seong Joon Oh, Sanghyuk Chun, Junsuk Choe, and Youngjoon Yoo. Cutmix: Regularization strategy to train strong classifiers with localizable features. In *IEEE/CVF International Conference on Computer Vision*, pages 6023–6032, 2019. doi: 10.1109/ICCV.2019.00612.
- [28] Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontiveros, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, and Amr Ahmed. Big bird: Transformers for longer sequences. *Advances in Neural Information Processing Systems*, 2020.
- [29] Hongyi Zhang, Moustapha Cisse, Yann N. Dauphin, and David Lopez-Paz. mixup: Beyond empirical risk minimization. In *International Conference on Learning Representations*, 2018.

A Full Woodbury Derivation

For $\gamma > 0$, using the Woodbury identity with base matrix $A_0 = \mu I_N$, low-rank factors $P = U$ and $Q = U^\top$, and middle matrix $M = \gamma I_r$ gives

$$(\mu I + \gamma U U^\top)^{-1} = \mu^{-1} I - \mu^{-1} U (\gamma^{-1} I + \mu^{-1} U^\top U)^{-1} U^\top \mu^{-1}.$$

Multiplying the small inverse by γ avoids an explicit γ^{-1} and gives the implementation form

$$Y = \mu^{-1} C - \mu^{-1} U (I + \gamma \mu^{-1} U^\top U)^{-1} (\gamma \mu^{-1} U^\top C),$$

which remains valid at $\gamma = 0$.